

# BUSA8001 (S2, 2023) Group Assignment - Predicting Used Car Sale Prices

---

**Kaggle Competition Ends:** Friday, 3 November 2023 @ 3:00pm (Week 13)

**Assignment Due Date on iLearn:** Friday, 3 November 2023 @ 11.59pm (Week 13)

## Overview:

- In the group assignment you will form a team of 3 students and participate in a forecasting competition on Kaggle
- The goal is to predict prices of used cars based on car characteristics and regression models
- Assessment Summary:
  - Write a problem statement and perform Exploratory Data Analysis
  - Clean up data, deal with categorical features and missing observations, and create new explanatory variables (feature engineering)
  - Construct and tune forecasting models, produce forecasts and submit your predictions to Kaggle
  - Each member of the team will record a video presentation of their work

## Instructions:

- Form a team of 3 students
- Each team member needs to join <https://www.kaggle.com>
- Choose a team leader and form a team in the competition  
<https://www.kaggle.com/t/32b34f072642495487836cf93453ac6a>
  - Team leader to click on `team` and join and invite other 2 team members to join
  - Your **team's name must start** with our unit code, for instance you could have a team called BUSA8001\_masterful\_geniuses
- All team members should work on all the tasks however
  - Choose a team member who will be responsible for one of each of the 3 tasks listed below
- Your predictions must be generated by a model you develop here
  - You will receive a mark of **zero** if your code does not produce the forecasts you submit to Kaggle

## Marks:

- Total Marks: 40
- Your individual mark will consist of:
  - $50\% \times \text{overall assignment mark} + 45\% \times \text{mark for the task that you are responsible for}$   
+  $5\% \times \text{mark received from your teammates for your effort in group work}$
- 1 mark: Ranking in the top 5 places of your unit on Kaggle
- 3 marks: Reaching the first place in the competition

**Submissions:**

1. On Kaggle: submit your team's forecast in order to be ranked by Kaggle
    - Limit of 20 submission per day
  2. On iLearn **only team leader to submit** this Jupyter notebook re-named `Group_Assignment_Team_Name.ipynb` where Team\_Name is your team's name on Kaggle
    - The Jupyter notebook must contain team members names/ID numbers, and team name in the competition
    - Provide answers to the 3 Tasks below in the allocated cells including all codes/outputs/writeups
    - One 15 minute video recording of your work
      - Each team member to provide a 5 minute presentation of the Task that they led (it is best to jointly record your video using Zoom)
      - When recording your video make sure your face is visible, that you share your Jupyter Notebook and explain everything you've done in the submitted Jupyter notebook on screen
      - 5 marks will be deducted from each Task for which there is no video presentation or if you don't follow the above instructions
  3. On iLearn each student needs to submit a file with their teammates' names, ID number and a mark for their group effort (out of 100%)
- 

**Fill out the following information**

For each team member provide name, Student ID number and which task is performed below

- Team Name on Kaggle: `(insert here)`
  - Team Leader and Team Member 1: `(insert here)`
  - Team Member 2: `(insert here)`
  - Team Member 3: `(insert here)`
- 

## Task 1: Problem Description and Initial Data Analysis

1. Read the Competition Overview on Kaggle  
<https://www.kaggle.com/t/32b34f072642495487836cf93453ac6a>
2. Referring to Competition Overview and the data provided on Kaggle write **Problem Description** (about 500 words) focusing on key points that will need to be addressed as first steps in Tasks 2 and 3 below,

- Using the following headings:
  - Forecasting Problem - explain what you are trying to do and how it could be used in the real world (i.e. why it may be important)
  - Evaluation Criteria - explain the criteria is used to assess forecast performance
  - Types of Variables/Features
  - Data summary and main data characteristics
  - Missing Values (only explain what you find - do not impute missing values at this stage)
  - You should **not** discuss any specific predictive algorithms at this stage
  - Note: Your written portion of this task should be completed in a single Markdown cell

Total Marks: 12

```
In [10]: # Importing necessary Libraries
import pandas as pd
```

```
# 1. Load the Data
filepath = "Users/user/KAGGLE"
testdf = pd.read_csv(filepath + "/test.csv")
traindf = pd.read_csv(filepath + "/train.csv")
```

```
In [2]: testdf.head()
```

|   | vin               | back_legroom | body_type       | city             | city_fuel_economy | daysonmarket | de... |
|---|-------------------|--------------|-----------------|------------------|-------------------|--------------|-------|
| 0 | 5N1AT2MV8HC824461 | 37.9 in      | SUV / Crossover | Wallingford      | 25.0              | 19           |       |
| 1 | 1GNSKCKC2HR160472 | 39 in        | SUV / Crossover | North Plainfield | 16.0              | 34           |       |
| 2 | 5NPD84LFXJH361029 | 35.7 in      | Sedan           | West Nyack       | 28.0              | 49           |       |
| 3 | 5XXGT4L34KG284916 | 35.6 in      | Sedan           | Bronx            | 24.0              | 20           |       |
| 4 | 2G1WF52E819291457 | NaN          | Sedan           | Little Ferry     | 19.0              | 38           |       |

5 rows × 38 columns

```
In [3]: traindf.head()
```

Out[3]:

|   | vin               | back_legroom | body_type       | city         | city_fuel_economy | daysonmarket | dealer_z |
|---|-------------------|--------------|-----------------|--------------|-------------------|--------------|----------|
| 0 | SJKCH5CRXHA032566 | 33.5 in      | SUV / Crossover | Great Neck   | NaN               | 20           | 1102     |
| 1 | 5LMCJ3D96HUL54638 | 36.8 in      | SUV / Crossover | Wayne        | 19.0              | 64           | 747      |
| 2 | 5LMCJ2D95HUL35217 | 36.8 in      | SUV / Crossover | Wayne        | 19.0              | 14           | 747      |
| 3 | 2HGFG1B86AH500600 | 30.3 in      | Coupe           | Little Ferry | 25.0              | 13           | 764      |
| 4 | 5LMCJ1D95LUL25032 | 38.6 in      | SUV / Crossover | Wayne        | 21.0              | 14           | 747      |

5 rows × 39 columns

```
In [4]: #Types of Data within dataset  
data_types = testdf.dtypes  
  
data_types
```

```
Out[4]:    vin          object
   back_legroom      object
   body_type         object
   city              object
   city_fuel_economy float64
   daysonmarket      int64
   dealer_zip        int64
   engine_displacement float64
   engine_type       object
   exterior_color    object
   franchise_dealer  bool
   front_legroom     object
   fuel_tank_volume  object
   fuel_type          object
   height             object
   highway_fuel_economy float64
   horsepower         float64
   interior_color    object
   is_new             bool
   latitude           float64
   length              object
   listed_date        object
   listing_color      object
   longitude           float64
   make_name          object
   maximum_seating    object
   mileage             float64
   model_name          object
   power               object
   savings_amount     int64
   seller_rating       float64
   torque              object
   transmission        object
   transmission_display object
   wheel_system        object
   wheelbase            object
   width                object
   year                 int64
dtype: object
```

```
In [5]: #Types of Data within dataset
data_types = traindf.dtypes
```

```
data_types
```

```
Out[5]:    vin          object
   back_legroom      object
   body_type         object
   city              object
   city_fuel_economy float64
   daysonmarket      int64
   dealer_zip        int64
   engine_displacement float64
   engine_type       object
   exterior_color    object
   franchise_dealer  bool
   front_legroom     object
   fuel_tank_volume  object
   fuel_type          object
   height             object
   highway_fuel_economy float64
   horsepower         float64
   interior_color    object
   is_new             bool
   latitude           float64
   length              object
   listed_date        object
   listing_color      object
   longitude           float64
   make_name          object
   maximum_seating    object
   mileage             float64
   model_name          object
   power               object
   savings_amount     int64
   seller_rating       float64
   torque              object
   transmission        object
   transmission_display object
   wheel_system        object
   wheelbase            object
   width                object
   year                 int64
   price                int64
   dtype: object
```

```
In [6]: # Data Summary and Exploration
summary_data_2 = traindf.describe(include='all')

summary_data_2
```

Out[6]:

|               | vin               | back_legroom | body_type       | city      | city_fuel_economy | daysonmarket | c           |
|---------------|-------------------|--------------|-----------------|-----------|-------------------|--------------|-------------|
| <b>count</b>  | 3500              | 3397         | 3494            | 3500      | 2912.000000       | 3500.000000  | 3500.000000 |
| <b>unique</b> | 3500              | 138          | 9               | 43        | NaN               | NaN          | NaN         |
| <b>top</b>    | SJKCH5CRXHA032566 | 38.2 in      | SUV / Crossover | Bay Shore | NaN               | NaN          | NaN         |
| <b>freq</b>   | 1                 | 207          | 1904            | 317       | NaN               | NaN          | NaN         |
| <b>mean</b>   | NaN               | NaN          | NaN             | NaN       | 21.655220         | 78.987714    | 118.0       |
| <b>std</b>    | NaN               | NaN          | NaN             | NaN       | 7.520448          | 104.838545   | 108.0       |
| <b>min</b>    | NaN               | NaN          | NaN             | NaN       | 10.000000         | 0.000000     | 9.0         |
| <b>25%</b>    | NaN               | NaN          | NaN             | NaN       | 18.000000         | 17.000000    | 70.0        |
| <b>50%</b>    | NaN               | NaN          | NaN             | NaN       | 21.000000         | 39.500000    | 88.0        |
| <b>75%</b>    | NaN               | NaN          | NaN             | NaN       | 24.000000         | 82.000000    | 117.0       |
| <b>max</b>    | NaN               | NaN          | NaN             | NaN       | 127.000000        | 698.000000   | 4920.0      |

11 rows × 39 columns

In [7]:

```
summary_data_2 = testdf.describe(include='all')

summary_data_2
```

Out[7]:

|               | vin               | back_legroom | body_type       | city      | city_fuel_economy | daysonmarket |     |
|---------------|-------------------|--------------|-----------------|-----------|-------------------|--------------|-----|
| <b>count</b>  | 1500              | 1447         | 1498            | 1500      | 1263.000000       | 1500.000000  | 15  |
| <b>unique</b> | 1500              | 117          | 9               | 43        | NaN               | NaN          | NaN |
| <b>top</b>    | 5N1AT2MV8HC824461 | 38.2 in      | SUV / Crossover | Bay Shore | NaN               | NaN          | NaN |
| <b>freq</b>   | 1                 | 89           | 795             | 141       | NaN               | NaN          | NaN |
| <b>mean</b>   | NaN               | NaN          | NaN             | NaN       | 21.239113         | 79.226667    | 124 |
| <b>std</b>    | NaN               | NaN          | NaN             | NaN       | 5.471276          | 115.406573   | 115 |
| <b>min</b>    | NaN               | NaN          | NaN             | NaN       | 12.000000         | 0.000000     | 5   |
| <b>25%</b>    | NaN               | NaN          | NaN             | NaN       | 18.000000         | 15.000000    | 70  |
| <b>50%</b>    | NaN               | NaN          | NaN             | NaN       | 21.000000         | 38.000000    | 102 |
| <b>75%</b>    | NaN               | NaN          | NaN             | NaN       | 24.000000         | 77.000000    | 117 |
| <b>max</b>    | NaN               | NaN          | NaN             | NaN       | 70.000000         | 1252.000000  | 492 |

11 rows × 38 columns

In [8]: # Check for Missing Values  
 missing\_values = traindf.isnull().sum()  
 missing\_values

```
Out[8]:    vin          0
            back_legroom   103
            body_type       6
            city           0
            city_fuel_economy 588
            daysonmarket     0
            dealer_zip       0
            engine_displacement 125
            engine_type      50
            exterior_color     0
            franchise_dealer   0
            front_legroom     103
            fuel_tank_volume   103
            fuel_type         37
            height           103
            highway_fuel_economy 588
            horsepower        125
            interior_color     0
            is_new            0
            latitude          0
            length            103
            listed_date       0
            listing_color      0
            longitude          0
            make_name          0
            maximum_seating    103
            mileage           203
            model_name         0
            power             299
            savings_amount      0
            seller_rating       0
            torque             331
            transmission        60
            transmission_display 60
            wheel_system       101
            wheelbase          103
            width              103
            year               0
            price              0
            dtype: int64
```

```
In [9]: # Check for Missing Values
missing_values = testdf.isnull().sum()

missing_values
```

```
Out[9]:    vin          0
           back_legroom   53
           body_type       2
           city            0
           city_fuel_economy 237
           daysonmarket     0
           dealer_zip        0
           engine_displacement 59
           engine_type       36
           exterior_color     0
           franchise_dealer    0
           front_legroom      53
           fuel_tank_volume    53
           fuel_type          29
           height            53
           highway_fuel_economy 237
           horsepower         59
           interior_color      0
           is_new             0
           latitude           0
           length             53
           listed_date        0
           listing_color       0
           longitude           0
           make_name          0
           maximum_seating     53
           mileage            76
           model_name          0
           power              126
           savings_amount      0
           seller_rating        3
           torque              141
           transmission         13
           transmission_display 13
           wheel_system         43
           wheelbase            53
           width               53
           year                0
           dtype: int64
```

```
In [10]: #shape of Train df
traindf.shape
```

```
Out[10]: (3500, 39)
```

```
In [11]: #shape of test df
testdf.shape
```

```
Out[11]: (1500, 38)
```

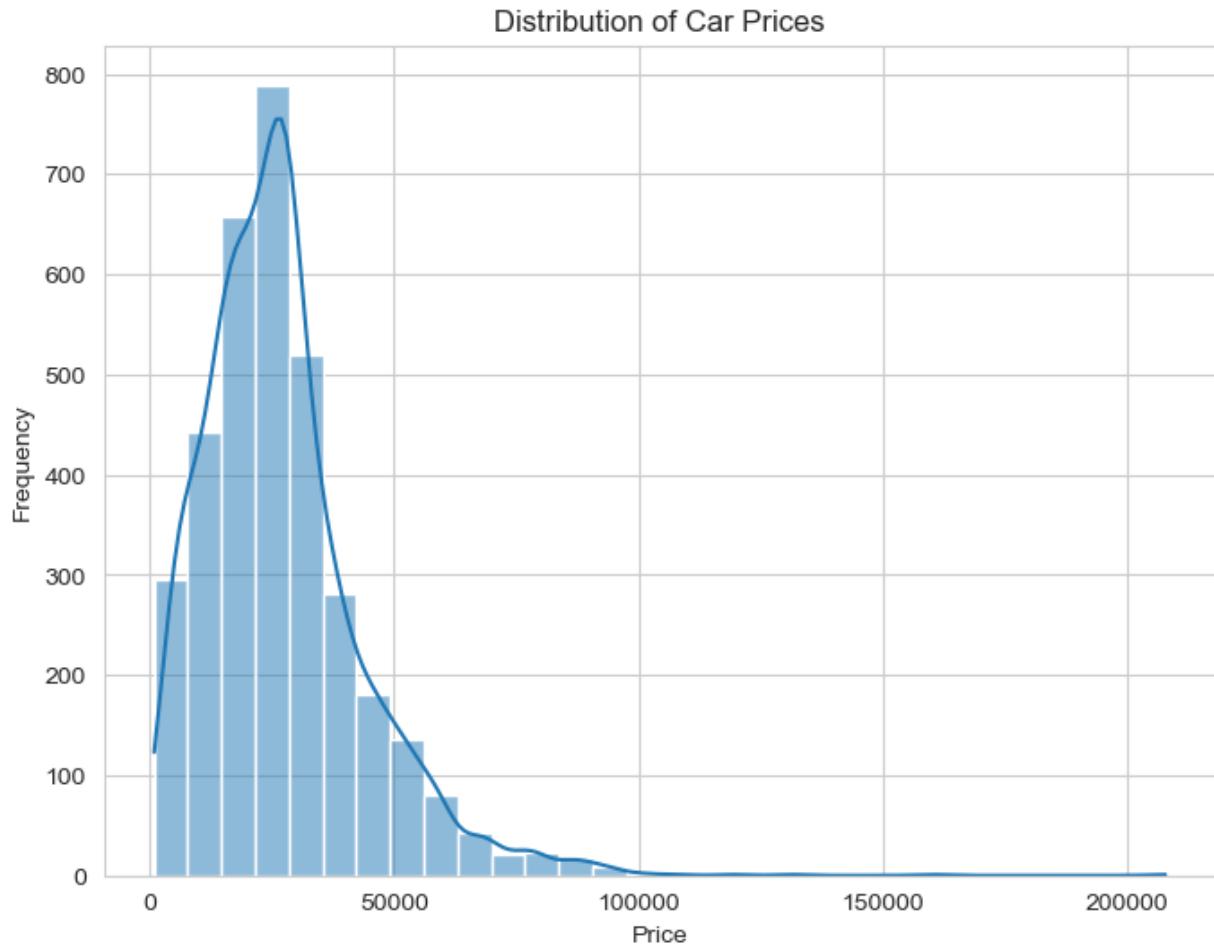
```
In [12]: #Exploratory Data Analysis based on train data
# Importing necessary Libraries for data visualization
import matplotlib.pyplot as plt
import seaborn as sns

sns.set_style("whitegrid")
plt.figure(figsize=(20, 15))
```

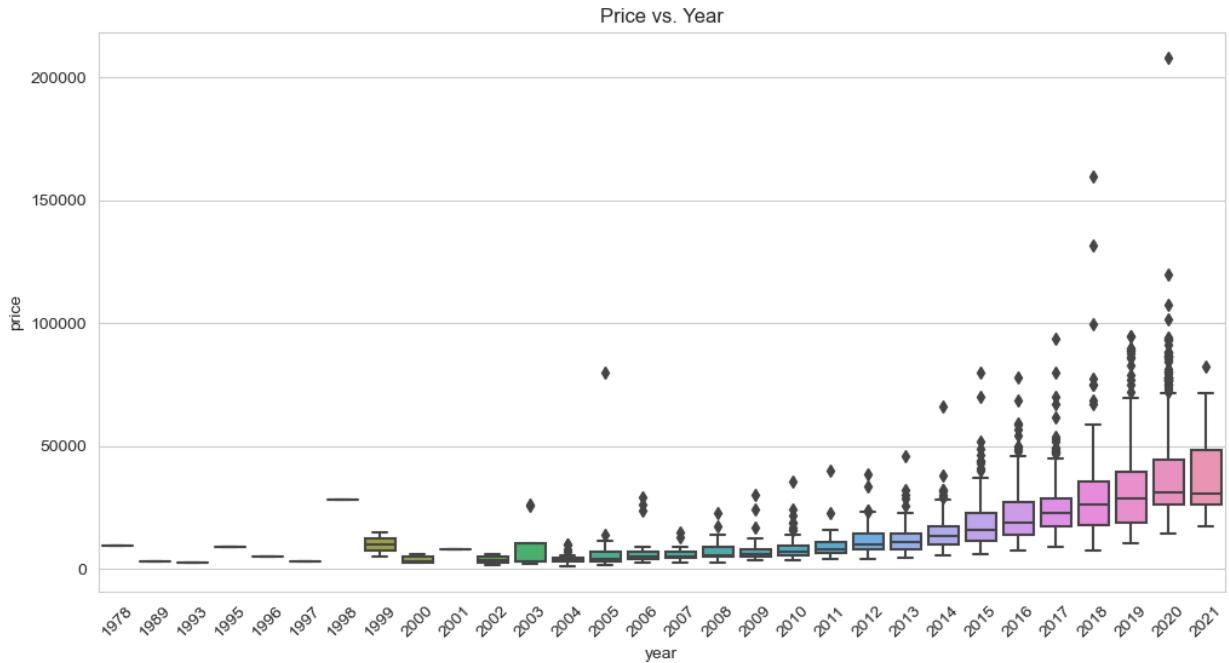
```
Out[12]: <Figure size 2000x1500 with 0 Axes>
```

```
<Figure size 2000x1500 with 0 Axes>
```

```
In [13]: # 1. Distribution of the Target Variable (price)
plt.figure(figsize=(8, 6))
sns.histplot(traindf['price'], bins=30, kde=True)
plt.title('Distribution of Car Prices')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()
```



```
In [14]: # 2. Price vs. Year
plt.figure(figsize=(12, 6))
sns.boxplot(x=traindf['year'], y=traindf['price'])
plt.title('Price vs. Year')
plt.xticks(rotation=45)
plt.show()
```

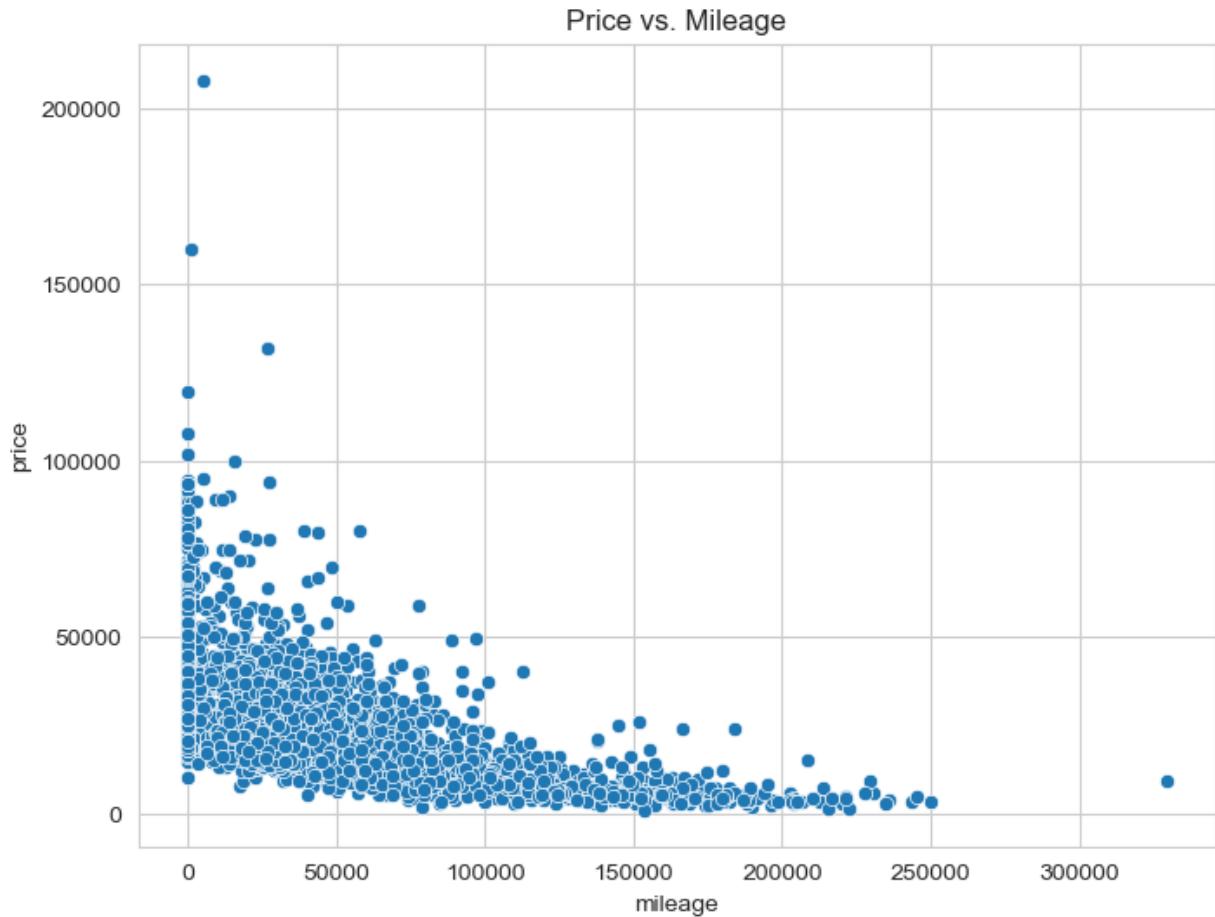


```
In [15]: # 3. Price vs. City Fuel Economy
plt.figure(figsize=(8, 6))
sns.scatterplot(x=traindf['city_fuel_economy'], y=traindf['price'])
plt.title('Price vs. City Fuel Economy')
plt.show()
```



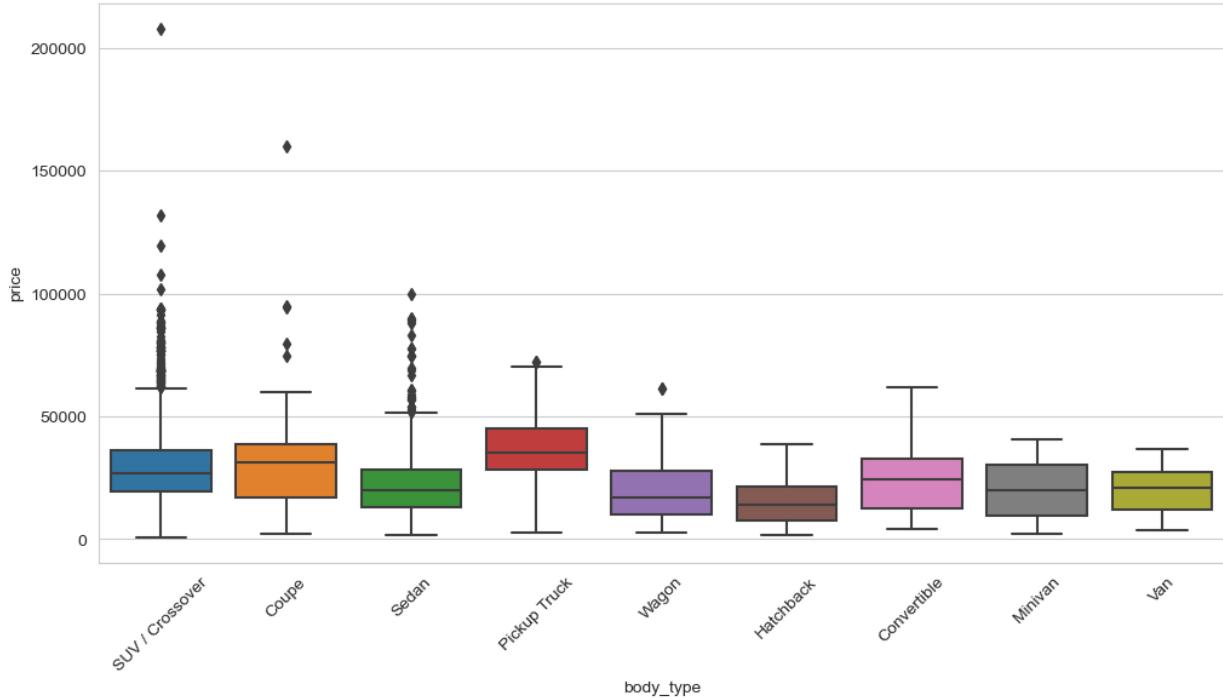
```
In [16]: # 4. Price vs. Mileage
plt.figure(figsize=(8, 6))
```

```
sns.scatterplot(x=traindf['mileage'], y=traindf['price'])
plt.title('Price vs. Mileage')
plt.show()
```



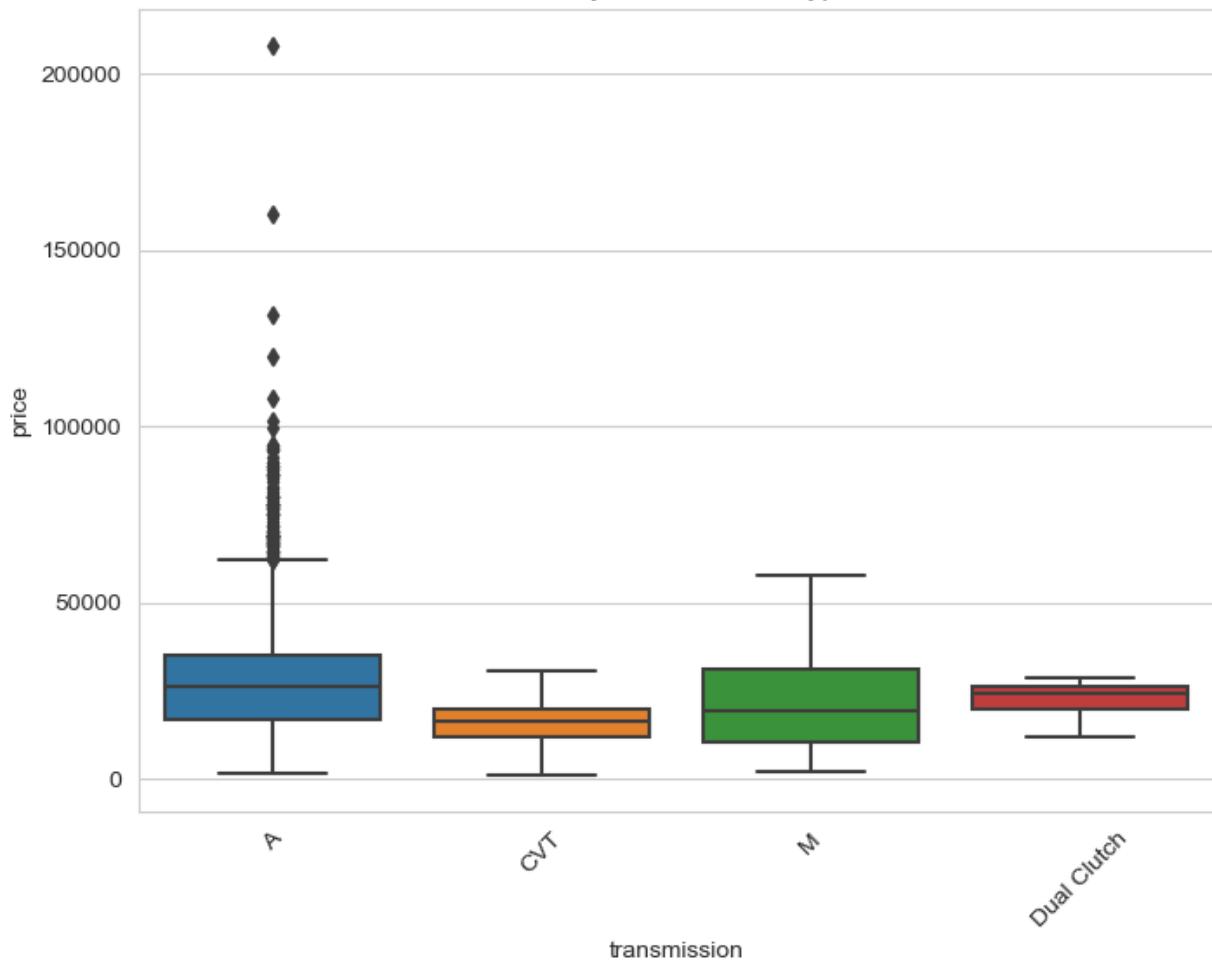
```
In [17]: # 5. Body Type Distribution
plt.figure(figsize=(12, 6))
sns.boxplot(x=traindf['body_type'], y=traindf['price'])
plt.title('Price Distribution by Body Type')
plt.xticks(rotation=45)
plt.show()
```

Price Distribution by Body Type



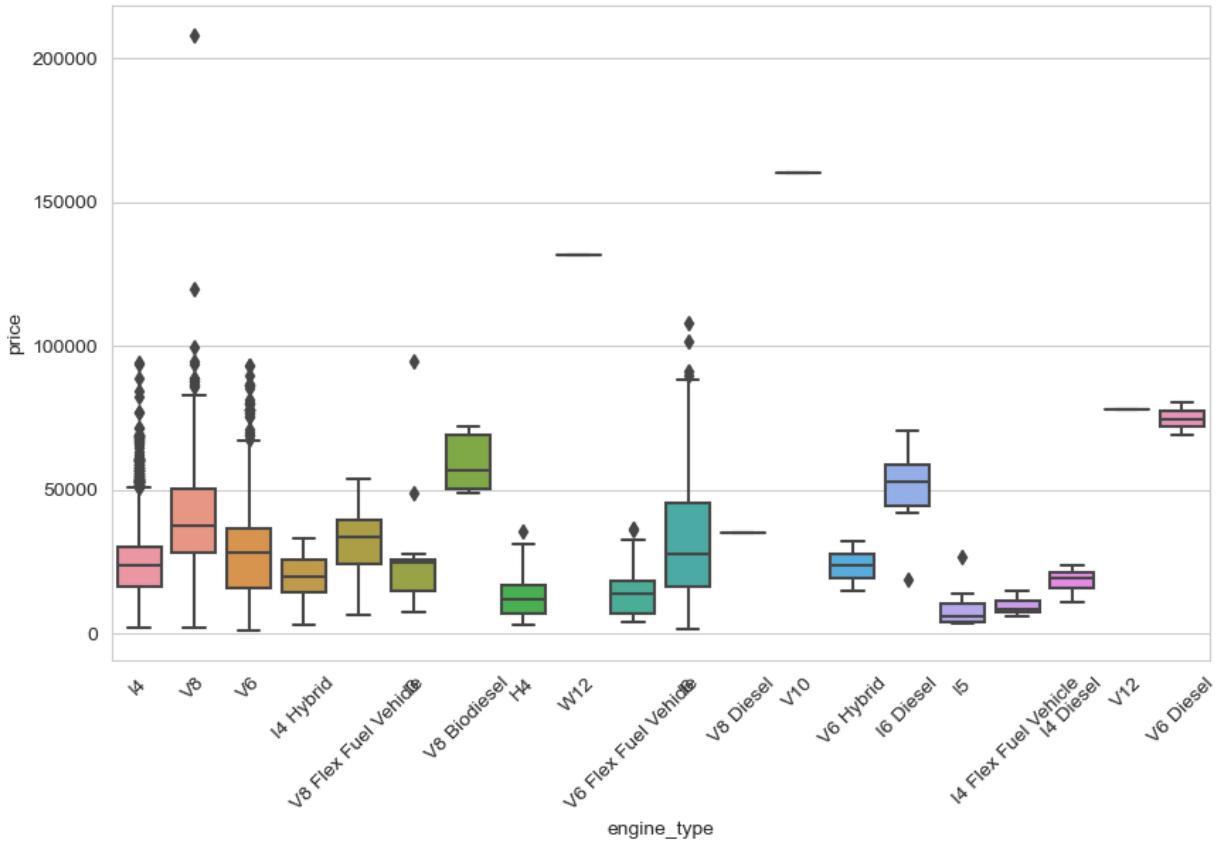
```
In [18]: # 6. Transmission Type vs. Price
plt.figure(figsize=(8, 6))
sns.boxplot(x=traindf['transmission'], y=traindf['price'])
plt.title('Price by Transmission Type')
plt.xticks(rotation=45)
plt.show()
```

## Price by Transmission Type



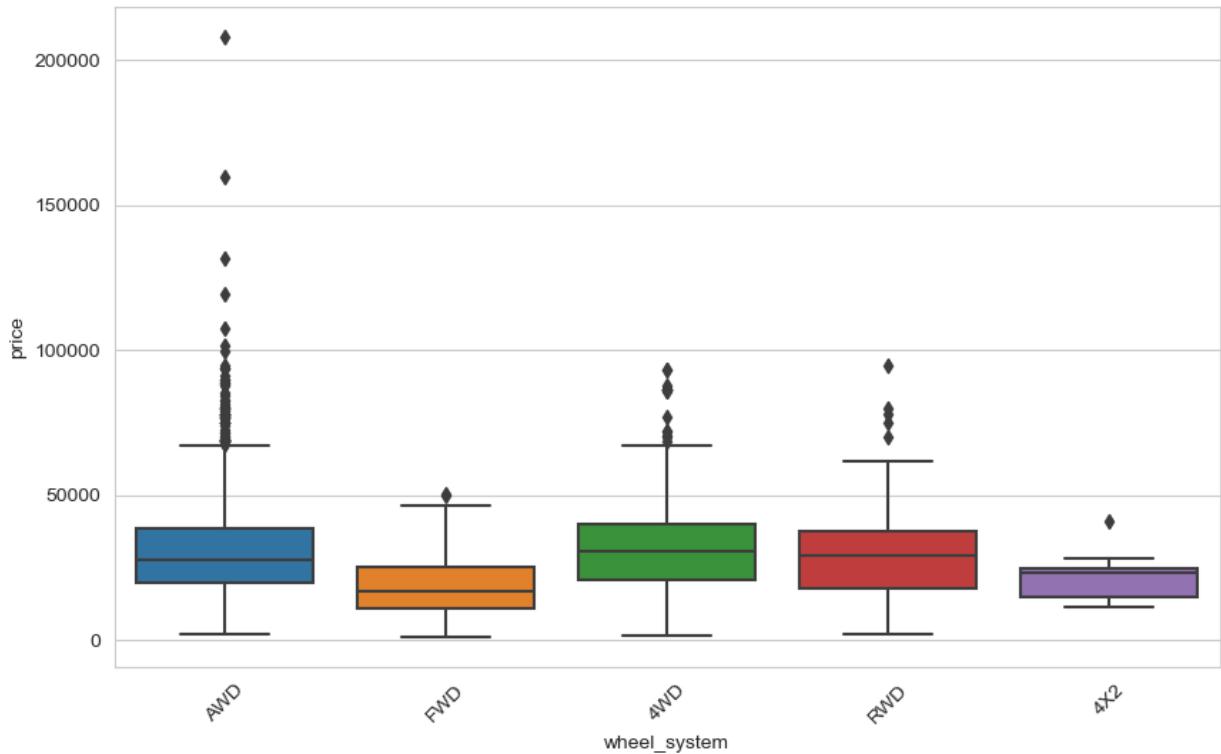
```
In [19]: # 7. Price vs. Engine Type
plt.figure(figsize=(10, 6))
sns.boxplot(x=traindf['engine_type'], y=traindf['price'])
plt.title('Price by Engine Type')
plt.xticks(rotation=45)
plt.show()
```

Price by Engine Type

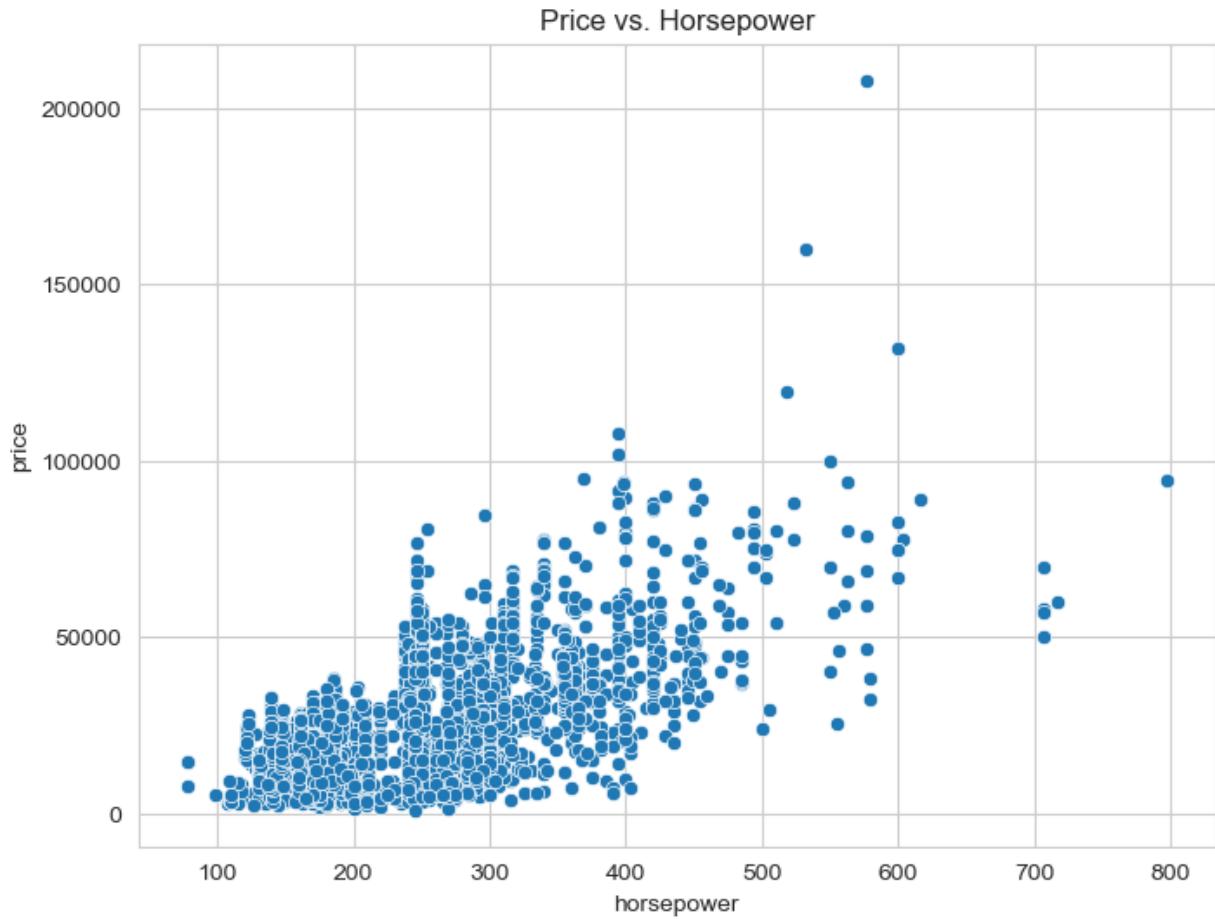


```
In [20]: # 8. Distribution of Wheel System
plt.figure(figsize=(10, 6))
sns.boxplot(x=traindf['wheel_system'], y=traindf['price'])
plt.title('Price by Wheel System')
plt.xticks(rotation=45)
plt.show()
```

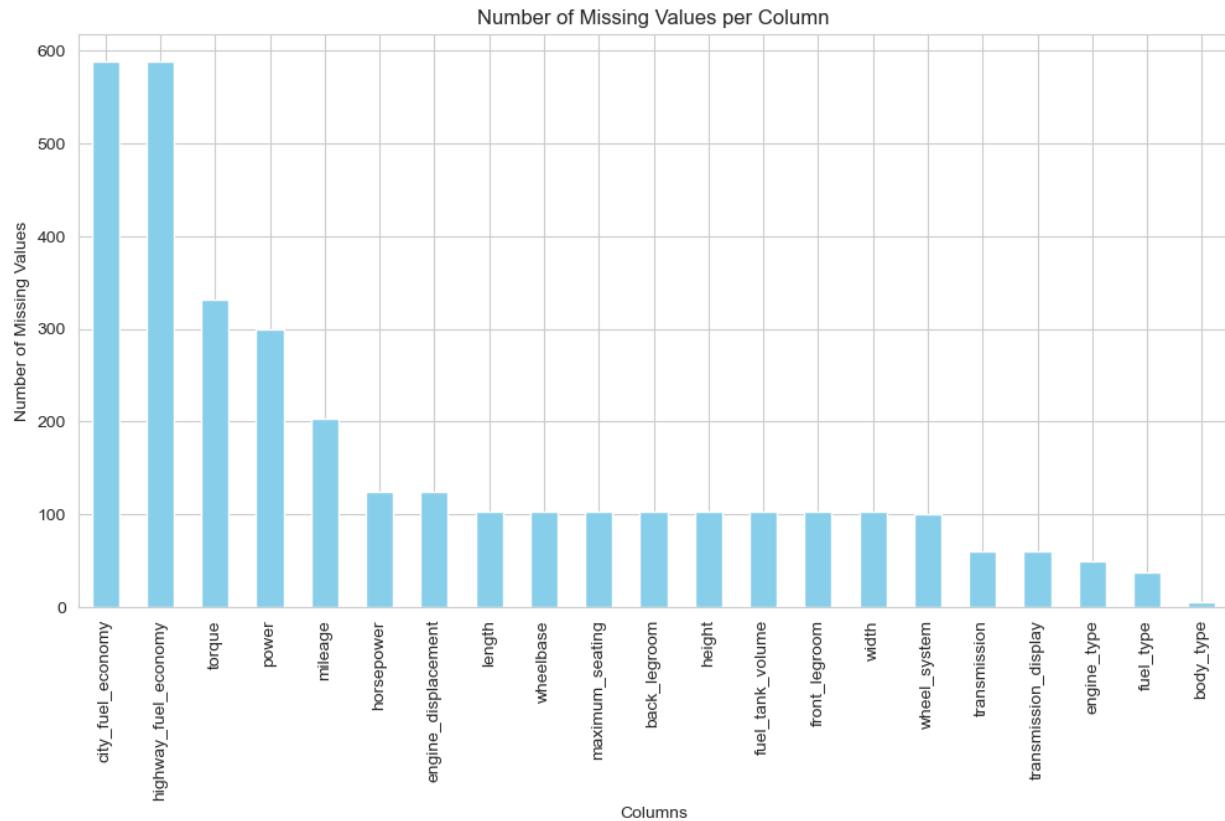
## Price by Wheel System



```
In [21]: # 9. Horsepower vs. Price
plt.figure(figsize=(8, 6))
sns.scatterplot(x=traindf['horsepower'], y=traindf['price'])
plt.title('Price vs. Horsepower')
plt.show()
```



```
In [22]: # 10. Number of Missing Values per Column
plt.figure(figsize=(12, 6))
missing_values_count = traindf.isnull().sum()
missing_values_count = missing_values_count[missing_values_count > 0].sort_values(ascending=True)
missing_values_count.plot(kind='bar', color='skyblue', title="Number of Missing Values")
plt.ylabel('Number of Missing Values')
plt.xlabel('Columns')
plt.show()
```



## Forecasting Problem

The objective is to forecast the cost of a used vehicle given its characteristics. To this end, we want to use data to develop a model that can predict a car's market worth based on a variety of criteria, including but not limited to its age, mileage, and engine type. In practical applications, this forecast is priceless. It guarantees a fair and lucrative pricing for the vendor. Instead, buyers may learn the vehicle's true market worth and avoid spending too much. The projections can also be useful for financial firms when determining loan amounts or insurance premiums. Thus, reliable pricing predictions improve honesty and confidence in the pre-owned vehicle trade.

## Evaluation Criteria

Metrics that judge how well and how accurately forecasts were made are the backbone of any forecasting competition. Although traditional regression measures are not explicitly stated in the published requirements, the focus on regression approaches like random forest and gradient boosting implies that they may be used in this competition. Metrics like Mean Absolute Error (MAE) and Root Mean Squared Error (RMSE), which penalise greater deviations more severely, and R-squared, which represents the fraction of variance explained by the model, are frequently used for regression issues. It's also possible to evaluate the model based on how well it deals with missing data, how well it incorporates engineering into its design, and how resilient it is in general.

## Types of Variables/Features

The dataset comprises a diverse set of variables that can be broadly categorized into the following types:

- Categorical Variables: These are variables that can take on one of a limited, and usually fixed, number of possible values. They represent qualitative data and are often non-numeric.
  - Examples from the dataset:
    - Sex: This can take values like 'Male' or 'Female'.
    - Marital status, Occupation, and Settlement size are other categorical variables in our dataset.
- Ordinal Categorical Variables: These are similar to categorical variables but have a clear order. They represent qualitative data where the order matters but the distance between categories is not defined.
  - Example from the dataset:
    - Education: It can have values like 'High School', 'Bachelor's', 'Master's', etc., where the order of education level is evident.
- Numerical Variables: These are variables that represent quantitative data and are numeric in nature. They can be further divided into:
  - Discrete: Variables that have a countable number of values. For instance, the number of doors in a car.
  - Continuous: Variables that can take any value within a range. For instance, age, income, or mileage of a car.
  - Examples from the dataset:
    - Age and Income are continuous numerical variables.
    - ID is a discrete numerical variable, representing unique identifiers for the entries.

Understanding the nature of these variables is fundamental as it guides the preprocessing steps, informs the choice of statistical methods, and impacts the model's performance. Different types of variables often require different handling techniques, especially in the realms of data visualization, feature engineering, and modeling.

## Data Summary and Main Data Characteristics

The dataset, train.csv, comprises 3,500 entries, each representing a unique car listing. Across these listings, there are 39 variables capturing diverse attributes of the cars, ranging from technical specifications to broader categorizations. The main objective of this dataset is to predict the target variable, price, based on the other 38 feature variables.

### Key Characteristics

- Diversity of Features: The dataset encapsulates a wide range of features, including categorical, ordinal, and numerical variables. This spectrum encompasses basic attributes like year and mileage to more nuanced ones such as body\_type and engine\_type.

- Distribution of Target Variable (price): Preliminary analysis indicates a right-skewed distribution, suggesting that a majority of cars are priced in the lower to mid-range bracket.
- Temporal Factor (year): The dataset spans various manufacturing years, potentially illustrating the depreciation of cars over time.
- Fuel Efficiency: Variables like city\_fuel\_economy and highway\_fuel\_economy provide insights into the car's fuel efficiency, a factor that might influence buyer decisions and consequently, the price.
- Variability in Car Types: The dataset is diverse in terms of car types, with multiple body\_type classifications like SUV, Sedan, and Truck. Such variety caters to different market segments and can significantly influence pricing.
- Technical Specifications: Features like horsepower, engine\_type, and transmission delve into the technical side of the vehicles, capturing performance and machinery aspects that can have direct implications on the car's value.
- Potential Data Quality Issues: While an in-depth cleaning process is yet to be undertaken, it's crucial to be vigilant about missing values, outliers, and potential inconsistencies in the dataset. An initial exploration did reveal some missing values, which will need addressing.

Understanding these characteristics is pivotal. It not only aids in formulating a suitable modeling strategy but also in conducting a comprehensive exploratory analysis, ensuring that the nuances of the dataset aren't overlooked.

## Missing Values

The train.csv dataset exhibits missing values across multiple columns. Key columns include city\_fuel\_economy, highway\_fuel\_economy, and various technical specifications like torque and power. Some columns, such as transmission and engine\_type, lack specific details about the car's machinery. Addressing these missing values is essential to ensure the robustness of predictive models.

### Methods to Handle Missing Values:

For continuous variables like city\_fuel\_economy, imputation with the median or mean is recommended. Technical specifications, being continuous, can be similarly treated. For categorical columns like transmission and engine\_type, the mode (most frequent value) is an appropriate imputation choice. In cases where relationships exist between variables, model-based imputation, using techniques like K-Nearest Neighbors or Random Forest, can be beneficial. For columns with a high proportion of missing values, consider dropping them. After imputation, validation through distribution checks is vital to ensure data integrity.

## Task 2: Data Cleaning, Missing Observations and Feature Engineering

- In this task you will follow a set of instructions/questions listed below.
- Make sure you explain each answer carefully both in Markdown text and on your video.

Total Marks: 12

**Task 2, Question 1:** Clean **all** numerical features so that they can be used in training algorithms.

For instance, back\_legroom feature is in object format containing both numerical values and text. Extract numerical values (equivalently eliminate the text) so that the numerical values can be used as a regular feature.

(2 marks)

```
In [2]: import pandas as pd
import re

# List of columns with mixed numeric and text values
columns_to_clean = ["back_legroom", "front_legroom", "fuel_tank_volume", "height", "length"]

# Function to extract numeric values and convert to float based on the format
def extract_numeric_value(text):
    if isinstance(text, str):
        numeric_part = re.search(r'\d+\.\d+', text)
        if numeric_part:
            return float(numeric_part.group())
    return None

for column in columns_to_clean:
    traindf[column] = traindf[column].apply(lambda x: extract_numeric_value(x))
    testdf[column] = testdf[column].apply(lambda x: extract_numeric_value(x))
```

I first identify the columns with mixed numeric and text values, and then create a function called extract\_numeric\_value. This function checks if a value is a text (string), and if so, it uses a regular expression to find and extract numeric patterns. The extracted numeric values are converted to floating-point numbers. If a valid numeric value is not found, the function returns None.

I then apply this extract\_numeric\_value function to clean the specified columns in both the traindf and testdf datasets. The result is that these columns will contain only the extracted numeric values, making them ready for analysis or modeling.

**Task 2, Question 2** Create at least 5 new features from the existing numerical variables which contain multiple items of information, for example you could extract maximum torque and torque rpm from the torque variable.

(2 marks)

```
In [3]: # Extract maximum torque from the 'torque' column
traindf['max_torque'] = traindf['torque'].str.extract(r'(\d+) lb-ft').astype(float)
testdf['max_torque'] = testdf['torque'].str.extract(r'(\d+) lb-ft').astype(float)
```

```

# Calculating the rear Legroom ratio
traindf['rear_legroom_ratio'] = traindf['back_legroom'] / traindf['front_legroom']
testdf['rear_legroom_ratio'] = testdf['back_legroom'] / testdf['front_legroom']

# Calculating fuel efficiency of the average of city and highway fuel economy
traindf['fuel_efficiency'] = (traindf['city_fuel_economy'] + traindf['highway_fuel_economy']) / 2
testdf['fuel_efficiency'] = (testdf['city_fuel_economy'] + testdf['highway_fuel_economy']) / 2

# Calculating the engine power ratio as horsepower / engine displacement
traindf['engine_power_ratio'] = traindf['horsepower'] / traindf['engine_displacement']
testdf['engine_power_ratio'] = testdf['horsepower'] / testdf['engine_displacement']

# Calculating car size as the product of length, width, and height
traindf['car_size'] = traindf['length'] * traindf['width'] * traindf['height']
testdf['car_size'] = testdf['length'] * testdf['width'] * testdf['height']

```

i created five features The maximum torque which indicates the maximum torque of cars, legroom ratio to show the ratio of rear legroom to front legroom, fuel efficiency of the average of city and highway fuel economy to reflects the average of city and highway fuel economy, engine power ratio as horsepower/engine displacement to denotes the ratio of horsepower to engine displacement, car sizes as product of length, width and height

**Task 2, Question 3:** Impute missing values for all features in both the training and test datasets. (3 marks)

```

In [5]: # Imputing missing values for numerical features with the median value
numerical_features = ["city_fuel_economy", "engine_displacement", "horsepower", "miles_per_gallon",
                      'car_age', 'max_torque', 'rear_legroom_ratio', 'fuel_efficiency',
                      "back_legroom", "front_legroom", "fuel_tank_volume", "height", "width",
                      'highway_fuel_economy']
for feature in numerical_features:
    traindf[feature].fillna(traindf[feature].median(), inplace=True)
    testdf[feature].fillna(testdf[feature].median(), inplace=True)

# Imputing missing values for categorical features with the mode (most frequent value)
categorical_features = ["body_type", "engine_type", "fuel_type", "transmission", "transverse_axis_transmission"]
for feature in categorical_features:
    traindf[feature].fillna(traindf[feature].mode()[0], inplace=True)
    testdf[feature].fillna(testdf[feature].mode()[0], inplace=True)

```

Missing values in numerical features are imputed with the median value, ensuring central tendencies. For categorical features, imputation with the mode (most frequent value) ensures representative values. This data preprocessing enhances dataset completeness for analysis and modeling

**Task 2, Question 4:** Encode all categorical variables appropriately as discussed in class.

- Where multiple values are given for an observation encode the observation as 'other'.
  - Where a categorical feature contains more than 5 unique values, map the features into 5 most frequent values + 'other' and then encode appropriately. For instance, map colours into 5 basic colours + 'other': [red, yellow, green, blue, purple, other] and then encode.
- (2 marks)

In [6]:

```
#custom mapping for 'listing_color'
listing_color_mapping = {
    'GRAY': 'GRAY',
    'BLACK': 'BLACK',
    'WHITE': 'WHITE',
    'SILVER': 'SILVER',
    'RED': 'RED'
}

# Encode 'listing_color'
traindf['listing_color'] = traindf['listing_color'].replace(listing_color_mapping)
testdf['listing_color'] = testdf['listing_color'].replace(listing_color_mapping)

other_colors = traindf['listing_color'].append(testdf['listing_color']).unique()
traindf['listing_color'] = traindf['listing_color'].apply(lambda x: x if x in other_colors else 'OTHERS')
testdf['listing_color'] = testdf['listing_color'].apply(lambda x: x if x in other_colors else 'OTHERS')

#custom mapping for 'maximum_seating'
seating_mapping = {
    '1 seats': '1 seats',
    '2 seats': '2 seats',
    '3 seats': '3 seats',
    '4 seats': '4 seats',
    '5 seats': '5 seats'
}

# Encode 'maximum_seating'
traindf['maximum_seating'] = traindf['maximum_seating'].replace(seating_mapping)
testdf['maximum_seating'] = testdf['maximum_seating'].replace(seating_mapping)

other_seating = traindf['maximum_seating'].append(testdf['maximum_seating']).unique()
traindf['maximum_seating'] = traindf['maximum_seating'].apply(lambda x: x if x in other_seating else 'OTHERS')
testdf['maximum_seating'] = testdf['maximum_seating'].apply(lambda x: x if x in other_seating else 'OTHERS')
```

C:\Users\user\AppData\Local\Temp\ipykernel\_6856\688642629.py:15: FutureWarning: The series.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

other\_colors = traindf['listing\_color'].append(testdf['listing\_color']).unique()

C:\Users\user\AppData\Local\Temp\ipykernel\_6856\688642629.py:33: FutureWarning: The series.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

other\_seating = traindf['maximum\_seating'].append(testdf['maximum\_seating']).unique()

I applied custom encoding to the 'listing\_color' variable, specifically mapping colors like gray, black, white, silver, and red to themselves, while categorizing all other colors as 'OTHERS.'

Similarly, for the 'maximum\_seating' variable, I encoded values ranging from 1 to 5 as they are, and categorized any values representing more seats as 'OTHERS.'

**Task 2, Question 5:** Perform any other actions you think need to be done on the data before constructing predictive models, and clearly explain what you have done.

(1 marks)

In [ ]:

```
# Calculating car age based on the 'year' column
traindf['car_age'] = 2023 - traindf['year']
testdf['car_age'] = 2023 - testdf['year']
```

I created a new column 'car age' and added it to the traindf and testdf I believe that the "car age" column is essential and plays a crucial role as it directly reflects the car's durability and longevity.

**Task 2, Question 6:** Perform exploratory data analysis to measure the relationship between the features and the target and carefully write up your findings. (2 marks)

In [11]:

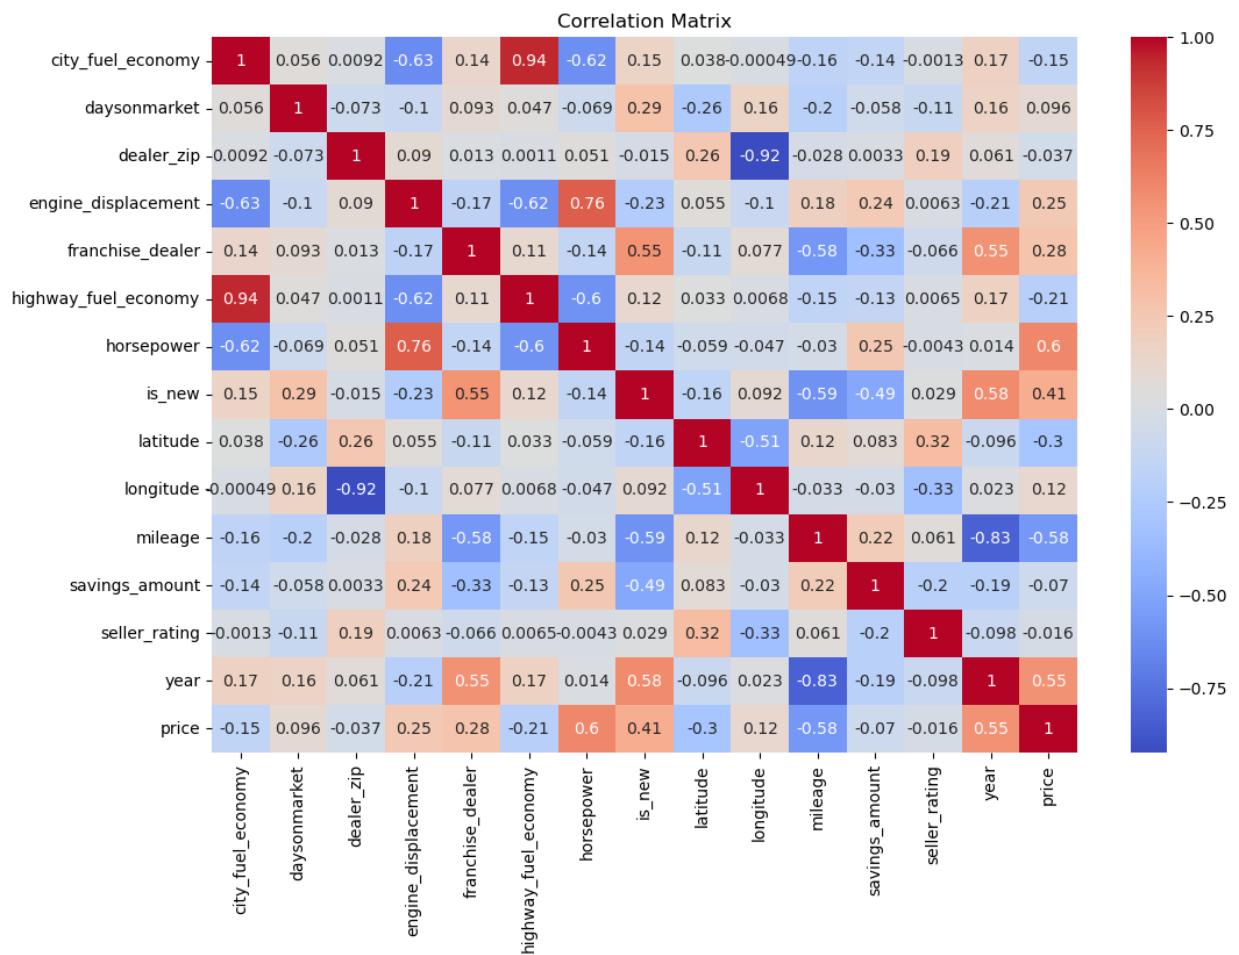
```

import seaborn as sns
import matplotlib.pyplot as plt

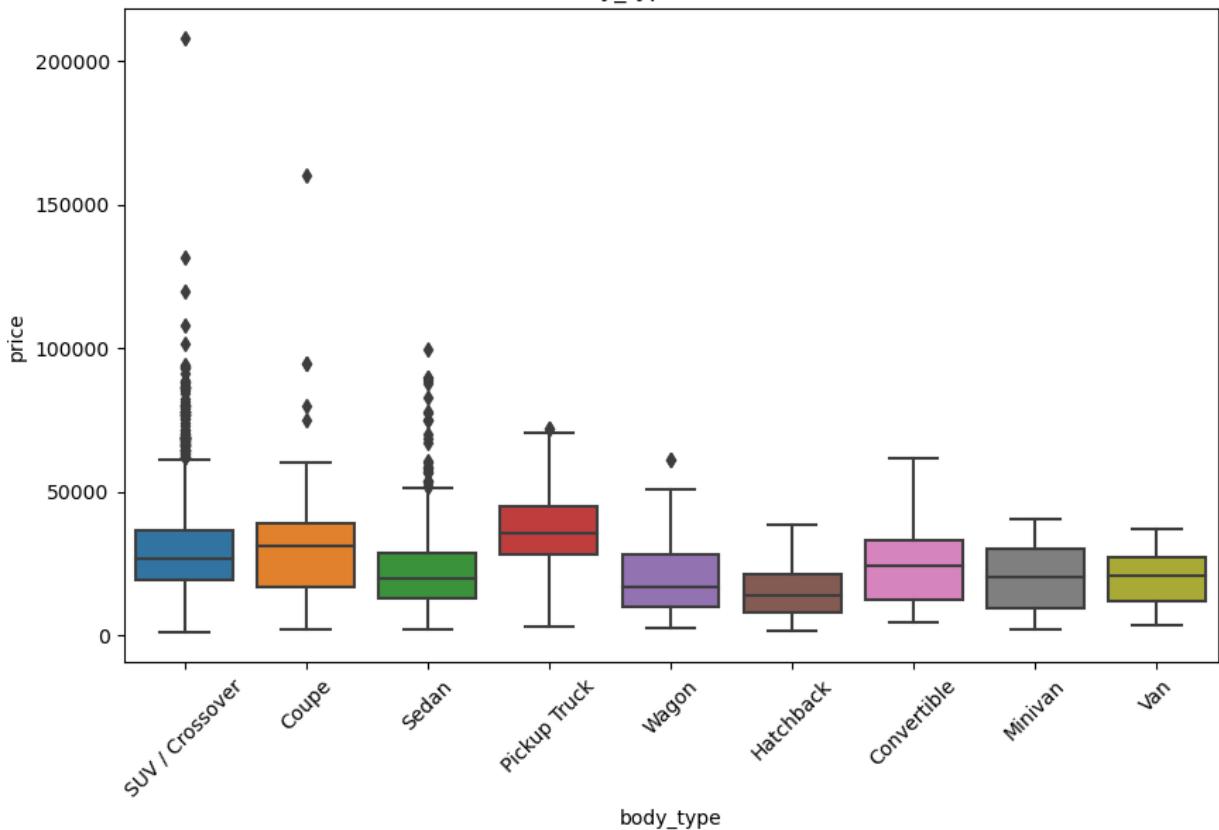
# Correlation matrix to measure the relationships between numerical features and the target
correlation_matrix = traindf.corr()
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()

# Boxplots for categorical features vs. target
categorical_features = ["body_type", "engine_type", "fuel_type", "transmission", "train"]
for feature in categorical_features:
    plt.figure(figsize=(10, 6))
    sns.boxplot(x=feature, y='price', data=traindf)
    plt.title(f'{feature} vs. Price')
    plt.xticks(rotation=45)
    plt.show()

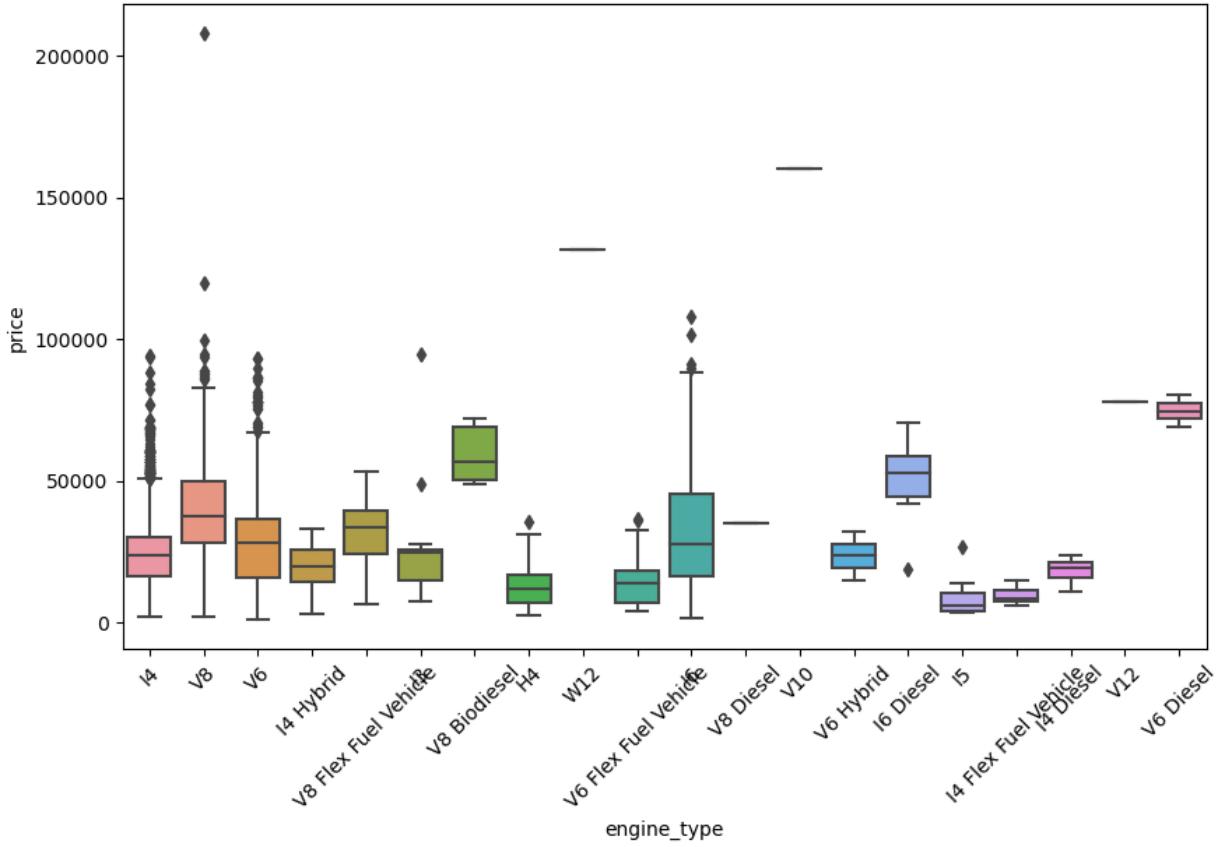
```



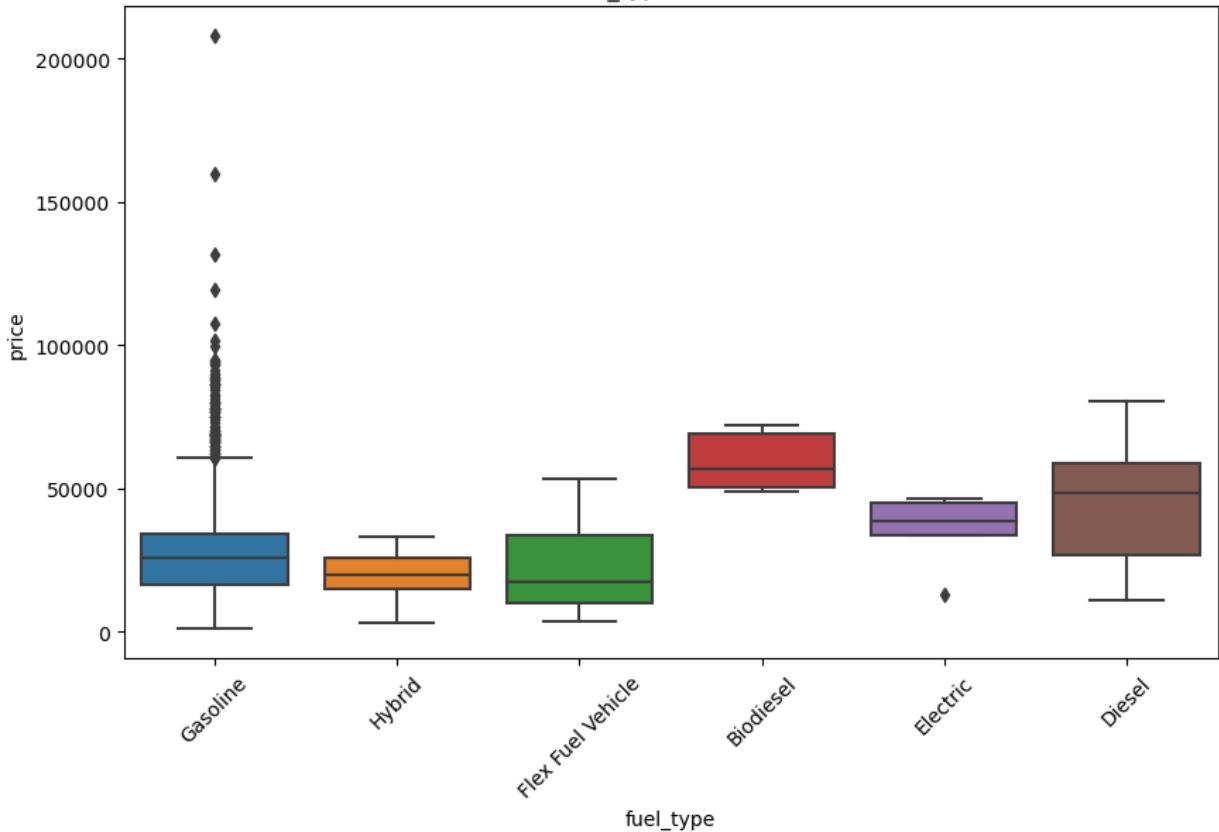
body\_type vs. Price



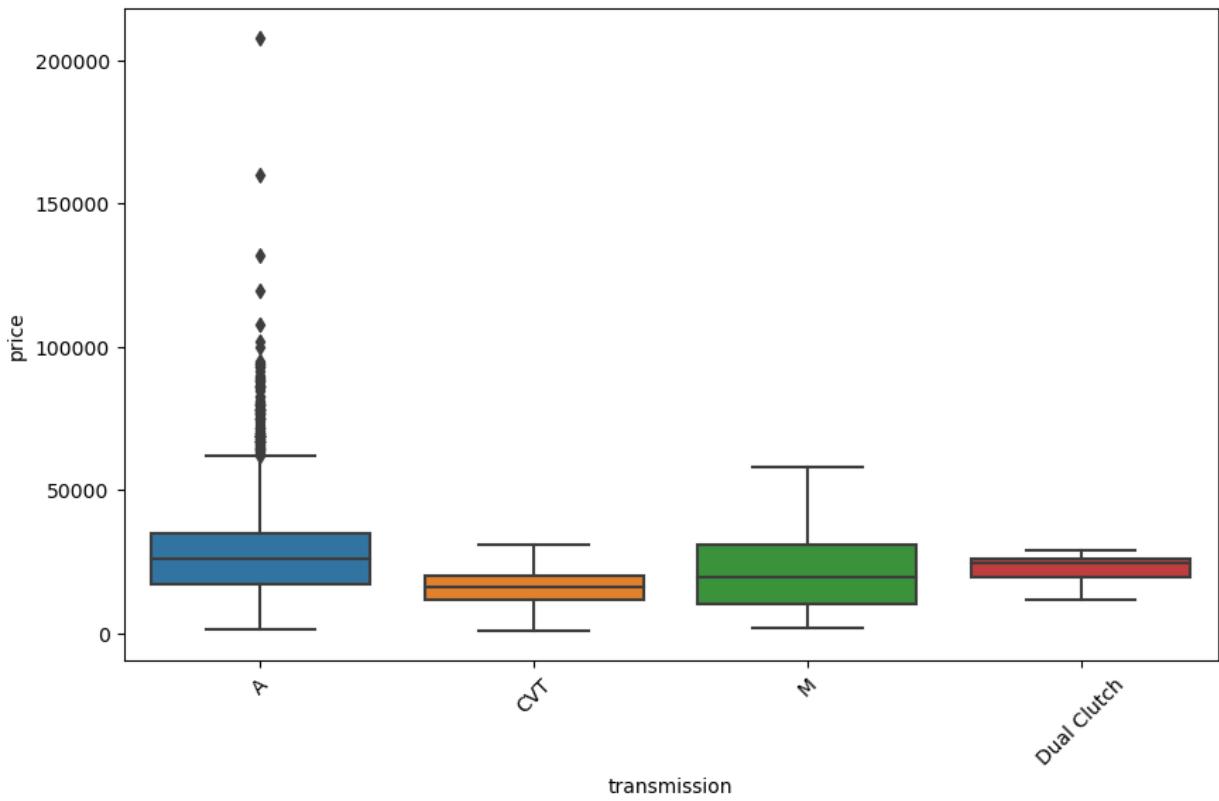
engine\_type vs. Price



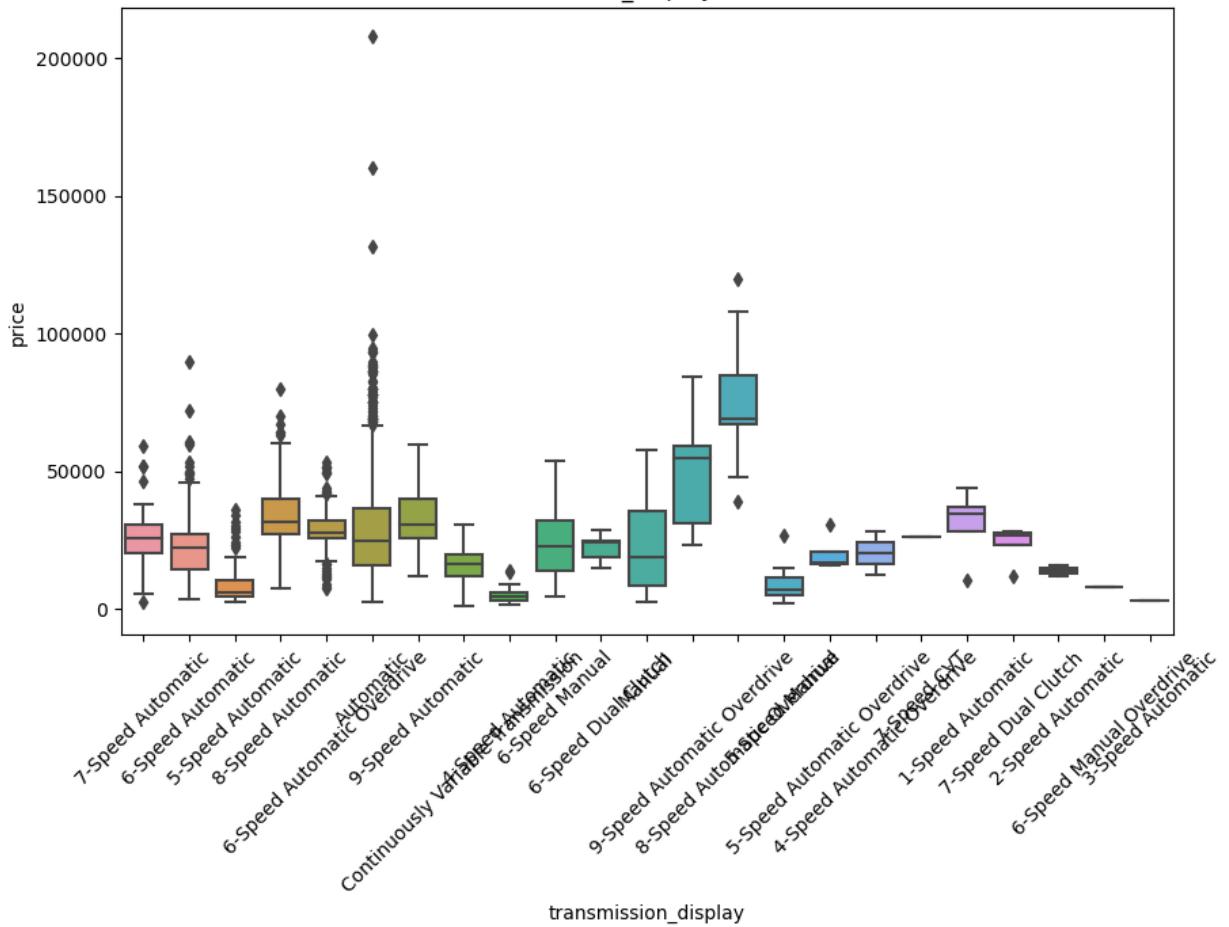
fuel\_type vs. Price



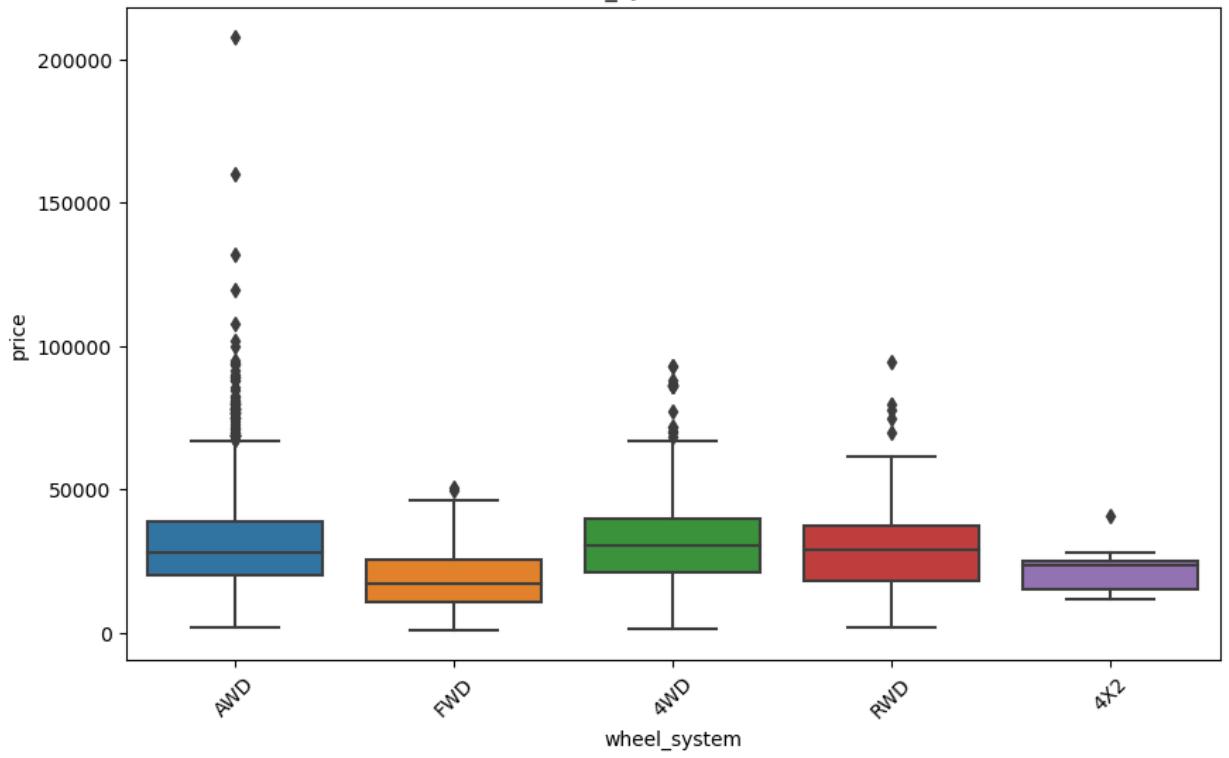
transmission vs. Price



## transmission\_display vs. Price



## wheel\_system vs. Price



The heatmap displays the relationships between variables using colors. The warmer colors like reds indicate correlations among all variables. In this case the variable "back\_legroom" stands out as having the correlation with the target variable 'price.'

Correlation Strength: When a variable has a correlation it means that as that variables value increases the cars 'price' tends to increase as well. On the hand if there is a negative correlation, an increase in that variables value is associated with a decrease, in the 'price.'

Positive and Negative Correlations: By looking at the heatmap we can determine which variables have correlations (represented by colors) and which ones have negative correlations (represented by cool colors) with 'price.'

Comparing Categorical Features to Price:

When examining the boxplots for features we gain extra insights, into how these features are connected to the target variable (price):

Body Type and Price: The boxplot for 'body\_type' indicates that pickup trucks generally have the highest average price, followed by coupes. However it's important to note that SUV/Crossover and Sedan categories have outliers suggesting that certain cars, in these categories have prices exceeding those of pickup trucks.

Engine. Price: By looking at the 'engine\_type' boxplot we can observe the price ranges associated with engine types. This information helps us understand what we can expect in terms of car prices. Notably some mid range cars equipped with V6 Diesel engines appear to have prices.

Fuel Type vs. Price: The 'fuel\_type' boxplot indicates that the most expensive cars are typically fueled by biodiesel. Gasoline-fueled cars are the most common, and there are outliers in various categories, suggesting that some cars in these categories can have notably high prices.

Transmission vs. Price: The 'transmission' boxplot shows that cars with 'A' (Automatic) transmission are the most common across all price ranges, including their outliers. Other transmission types appear to be more randomly scattered.

Transmission Display vs. Price: The 'transmission\_display' boxplot displays different types of transmission displays and their distribution in the dataset. Notably, cars with '9-speed Automatic Overdrive' have the highest price range, while those with '8-speed Automatic' are spread out and contain outliers.

Wheel System vs. Price: The 'wheel\_system' boxplot illustrates the distribution of various wheel systems in the dataset. 'AWD' (All-Wheel Drive) is the most common wheel system, followed by '4WD' (Four-Wheel Drive), and other categories.

---

## Task 3: Fit and tune a forecasting model/Submit predictions/Report score and ranking

Make sure you **clearly explain each step** you do both in text and on the recorded video.

This task must not create any additional features and has to use on the dataset constructed in

## Task 2.

1. Build at least 3 machine learning (ML) regression models taking into account the outcomes of Tasks 1 & 2 (Explain Carefully)
  2. Fit the models and tune hyperparameters via cross-validation: make sure you comment and explain each step clearly
  3. Select your best algorithm, create predictions using the test dataset, and submit your predictions on Kaggle's competition page
  4. Provide Kaggle ranking and **score** (screenshot your best submission) and comment
  5. Make sure your Python code works, so that a marker that can replicate your all Kaggle Score
- Hint: to perform well you will need to iterate Tasks 2 and Task 3.

Total Marks: 12

In [ ]: #Task 3 code here

(Task 3 - insert more cells as required)

## Marking Criteria

- Marking Rubrics
  - Problem Description - 12 marks
  - Data Cleaning - 12 marks
  - Building Forecasting models - 12 marks
  - Competition Points - 4 marks
- To receive full marks your solutions must satisfy the following criteria:
  - Provide Python solutions that follow the modelling methodology developed in BUSA8001
  - Written answers explain your logic and Python code in detail, and be formulated in easy to understand full sentences